

DAY 2 · MODULE 3 · 25 MIN

Evaluating Retrieval

Know if your knowledge layer actually works

Why evaluate retrieval separately

If you skip this step you'll do one of two things:

- Ship an agent that seems to work — until it doesn't, and they can't tell why
- Tune prompts forever, when the real problem was chunking or embeddings

Retrieval failures and generation failures look identical from the outside. Evaluate separately, fix separately.

Day 4 anchors evaluation for full workflows. Today's module gets the habit started.

Two evaluation modes for RAG

Mode	What it evaluates	Ground truth required
Process evaluation	The retrieval step itself — did we get relevant chunks?	Depends on evaluator
System evaluation	The end-to-end response — did the agent answer correctly?	Sometimes

- Process eval finds the retrieval-side bug
- System eval confirms the whole thing works
- You'll do a small dose of each in today's lab

The Foundry evaluator catalog for RAG

Evaluator	Type	Needs ground truth?
Retrieval	Process	No — LLM judges context relevance
Groundedness	System	No — LLM judges if response is grounded in context
Groundedness Pro (preview)	System	No — Content Safety service, boolean
Relevance	System	No — LLM judges if response addresses query
Response Completeness (preview)	System	Yes — needs expected answer
Document Retrieval	Process	Yes — computes metrics from qrels labels (no LLM judge)

Start with the top four. Add ground truth when you're ready to invest in labels.

Where to start — Retrieval and Groundedness

Two zero-setup RAG evaluators — one process, one system. Neither needs ground truth.

Retrieval (process)

- Input: query, context (query optional; both improve scoring)
- Output: 1–5 score (pass ≥ 3), plus reasoning
- Answers: 'Are the chunks we pulled actually relevant to the question?'

Groundedness (system)

- Input: response required; context recommended; query optional
- Output: 1–5 score (pass ≥ 3), plus reasoning
- Answers: 'Did the response stay in the context, or did the model fabricate?'

Together, these two isolate retrieval-side vs. generation-side failure.

D E M O

~4 min

Score two agents live

Two static transcripts — grounded (docs-assistant WITH IQ attached) and baseline (no knowledge source) — scored side-by-side with `retrieval_eval.py`. Grounded run clears the Retrieval ≥ 3.5 and Groundedness ≥ 4.0 thresholds. Baseline run refuses to score at all — the eval script hard-fails on empty context. That refusal is a real production signal.

Runbook: [demos/day2/module-3-demo-1-score-two-agents.md](#) · Files: [demos/day2/module-3-demo-1-score-two-agents/](#)

Groundedness vs. Response Completeness

Two sides of the same coin:

- Groundedness = precision. Did the response contain anything not in the context? (Fabrication check.)
- Response Completeness = recall. Did the response leave anything out that the ground truth expected? (Coverage check.)

A high-groundedness, low-completeness response is accurate but partial. A low-groundedness, high-completeness response is comprehensive but making things up. Track both.

Response Completeness needs ground truth — invest when you have a benchmark corpus.

Document Retrieval — the deep debug tool

When retrieval quality is the bottleneck, this is your parameter-sweep evaluator.

- Input: query-relevance labels (qrels) and the ranked retrieval output
- Output: NDCG, XDCG, Fidelity, Max Relevance, Holes at various k
- Use it to answer:
 - Should I use vector, semantic, or hybrid?
 - What's the right top_k?
 - Is 500-token chunking better than 1000?

You need labeled data — a person judged which docs are relevant for each query. Worth it when tuning matters.

Configuring an evaluator (Python SDK)

```
testing_criteria = [  
    {  
        "type": "azure_ai_evaluator",  
        "name": "groundedness",  
        "evaluator_name": "builtin.groundedness",  
        "initialization_parameters": {"deployment_name": model_deployment},  
        "data_mapping": {  
            "context": "{{item.context}}",  
            "response": "{{item.response}}",  
        },  
    },  
    {  
        "type": "azure_ai_evaluator",  
        "name": "retrieval",  
        "evaluator_name": "builtin.retrieval",  
        "initialization_parameters": {"deployment_name": model_deployment},  
        "data_mapping": {"query": "{{item.query}}", "context": "{{item.context}}"},  
    },  
]
```

Same JSON shape for each evaluator. The data_mapping tells the evaluator where to find fields on your test dataset.

The test dataset — smaller than you think

A JSONL file, one line per test case:

```
{ "query": "How do I set up a Prompt agent?",  
  "context": "A Prompt agent is authored in the Foundry portal or SDK...",  
  "response": "You create a Prompt agent by..." }  
{ "query": "What models does Foundry support?",  
  "context": "Foundry hosts models from Azure OpenAI, Anthropic, Meta...",  
  "response": "Foundry supports multiple models including..." }
```

- Start with 10–15 examples
- Enough signal to catch obvious regressions; small enough to actually maintain
- Add examples over time as you find failure modes

Reading the output

```
{
  "type": "azure_ai_evaluator",
  "name": "Groundedness",
  "metric": "groundedness",
  "score": 4,
  "label": "pass",
  "reason": "The response is well-grounded without fabricating content.",
  "threshold": 3,
  "passed": true
}
```

- score — 1–5 numeric (or true/false for Pro)
- label / passed — pass or fail against the threshold
- reason — the LLM judge's explanation (read this when things fail)
- metric — which evaluator produced this row

The reason field is where debugging happens. Read it.

The iteration loop

Same shape as Day 1's prompt-engineering loop:

- Author an eval dataset (10–15 items)
- Run Retrieval + Groundedness against your current pipeline
- Read the failures — retrieval bug or generation bug?
- Change one thing (chunk size, top_k, prompt, model)
- Rerun. Compare scores.

Change one thing at a time. Three changes at once = you learned nothing about what worked.

Same discipline shows up Day 4 for multi-agent workflow eval.

LLM-as-judge caveats

You're using an LLM to judge an LLM. Some caveats:

- Judge bias — the judge model has preferences. Same prompt, different judges = different scores.
- Judge cost — every eval item is an extra LLM call. Budget accordingly.
- Judge drift — model updates can shift baseline scores. Version-pin your judge.
- Judge as a check, not truth — cross-check with the reason field and spot-check the actual failures.

Day 5 covers online eval and drift more deeply. Today: know the judge isn't infallible.

Common traps

- Testing end-to-end only — a bad answer might be retrieval OR generation; can't tell without process eval
- Test set too small — 3 items = coin flip; 10–15 is the useful floor
- Test set too large — every eval run costs LLM calls; 50 hand-labeled > 5,000 unlabeled
- Changing three things at once — you learn nothing about what worked
- Ignoring the reason field — score tells you if it failed; reason tells you why
- Judge model = production model — conflict of interest; use a different (usually smaller) judge

Where eval shows up the rest of the week

- Day 3 (single-agent depth) — evaluator patterns extend to tool-use correctness
- Day 4 (multi-agent anchor) — trajectory eval, cost-per-successful-outcome, regression harness
- Day 5 (production) — online eval, drift detection, red-teaming
- Capstone — required elements include a golden set + eval scores

The habit that starts today runs through the whole week.

Takeaways

- Evaluate retrieval separately or you can't tell retrieval bugs from generation bugs.
- Foundry evaluators: start with Retrieval (process) + Groundedness (system) — both zero-setup.
- Groundedness (precision) + Response Completeness (recall) — track both when you have ground truth.
- Small test sets beat big untested claims — 10–15 hand-labeled items is the floor.
- Change one thing, then rerun — same iteration loop as prompt engineering.

Next: Tools layer — how agents do things beyond talking.